

Bank customer churn analysis

In []: In the rapidly evolving banking sector, customer retention has become a critical factor that influence customer decisions to stay **with** or leave their banking service. We will analyze various attributes of bank customers to identify key predictors of customer churn and gain insights that could help devise strategies to enhance customer retention **and** reduce churn.

```
In [1]: #Importing necessary libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [196... df = pd.read_csv('https://drive.google.com/uc?export=download&id=1xh7D0NDmxc')
df.head()
```

```
Out[196... 
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

```
In [11]: df.shape
```

```
Out[11]: (10000, 18)
```

```
In [ ]: # RowNumber, surname and CustomerId is irrelevant, lets delete it
```

```
In [197... df.drop(['RowNumber', 'CustomerId', 'Surname'], axis = 1, inplace = True)
df.head()
```

Out[197...

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts
0	619	France	Female	42	2	0.00	1
1	608	Spain	Female	41	1	83807.86	1
2	502	France	Female	42	8	159660.80	3
3	699	France	Female	39	1	0.00	2
4	850	Spain	Female	43	2	125510.82	1

In [13]: `df.columns`

Out[13]: Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited', 'Complain', 'Satisfaction Score', 'Card Type', 'Point Earned'], dtype='object')

In [14]: `df.dtypes`

Out[14]: CreditScore int64
Geography object
Gender object
Age int64
Tenure int64
Balance float64
NumOfProducts int64
HasCrCard int64
IsActiveMember int64
EstimatedSalary float64
Exited int64
Complain int64
Satisfaction Score int64
Card Type object
Point Earned int64
dtype: object

In [15]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CreditScore            10000 non-null  int64
1   Geography              10000 non-null  object
2   Gender                 10000 non-null  object
3   Age                   10000 non-null  int64
4   Tenure                 10000 non-null  int64
5   Balance                10000 non-null  float64
6   NumOfProducts          10000 non-null  int64
7   HasCrCard              10000 non-null  int64
8   IsActiveMember         10000 non-null  int64
9   EstimatedSalary        10000 non-null  float64
10  Exited                  10000 non-null  int64
11  Complain                10000 non-null  int64
12  Satisfaction Score     10000 non-null  int64
13  Card Type               10000 non-null  object
14  Point Earned            10000 non-null  int64
dtypes: float64(2), int64(10), object(3)
memory usage: 1.1+ MB

```

```
In [16]: df.isna().sum()
```

```

Out[16]: CreditScore      0
         Geography        0
         Gender           0
         Age              0
         Tenure            0
         Balance           0
         NumOfProducts     0
         HasCrCard         0
         IsActiveMember    0
         EstimatedSalary    0
         Exited            0
         Complain          0
         Satisfaction Score 0
         Card Type         0
         Point Earned      0
         dtype: int64

```

```
In [17]: df.duplicated().sum()
```

```
Out[17]: 0
```

```

In [198... df['HasCrCard'].replace({0 : 'No', 1 : 'Yes'}, inplace = True)
df['IsActiveMember'].replace({0 : 'No', 1 : 'Yes'}, inplace = True)
df['Exited'].replace({0 : 'No', 1 : 'Yes'}, inplace = True)
df['Complain'].replace({0 : 'No', 1 : 'Yes'}, inplace = True)
df.head()

```

Out[198...

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts
0	619	France	Female	42	2	0.00	1
1	608	Spain	Female	41	1	83807.86	1
2	502	France	Female	42	8	159660.80	3
3	699	France	Female	39	1	0.00	2
4	850	Spain	Female	43	2	125510.82	1

In [200...

```
columnss = ['Geography', 'Gender', 'NumOfProducts', 'HasCrCard', 'IsActiveMe  
for i in columnss:  
    print(df.value_counts(i))  
    print('\n')
```

Geography
France 5014
Germany 2509
Spain 2477
Name: count, dtype: int64

Gender
Male 5457
Female 4543
Name: count, dtype: int64

NumOfProducts
1 5084
2 4590
3 266
4 60
Name: count, dtype: int64

HasCrCard
Yes 7055
No 2945
Name: count, dtype: int64

IsActiveMember
Yes 5151
No 4849
Name: count, dtype: int64

Exited
No 7962
Yes 2038
Name: count, dtype: int64


Complain
No 7956
Yes 2044
Name: count, dtype: int64

Satisfaction Score
3 2042
2 2014
4 2008
5 2004
1 1932
Name: count, dtype: int64

Card Type
DIAMOND 2507

```
GOLD      2502
SILVER    2496
PLATINUM  2495
Name: count, dtype: int64
```

```
In [133... # Check proportion of customer exited
data['Exited'].value_counts().plot.pie(autopct = '%.1f%%', explode =
(0,0.1))
plt.title('Proportion of Customer Exited')
plt.show()
```



20.4 % of the customer have exited the bank

Exploratory Data Analysis (EDA)

Statistical Summary

```
In [143... #Non Graphical Analysis

columnss = ['Geography', 'Gender', 'NumOfProducts', 'HasCrCard', 'IsActiveMe
for i in columnss:
    print(df.value_counts(i))
    print('\n')
```

Geography
0 5014
1 2509
2 2477
Name: count, dtype: int64

Gender
1 5457
0 4543
Name: count, dtype: int64

NumOfProducts
1 5084
2 4590
3 266
4 60
Name: count, dtype: int64

HasCrCard
1 7055
0 2945
Name: count, dtype: int64

IsActiveMember
1 5151
0 4849
Name: count, dtype: int64

Exited
0 7962
1 2038
Name: count, dtype: int64

Complain
0 7956
1 2044
Name: count, dtype: int64

Satisfaction Score
3 2042
2 2014
4 2008
5 2004
1 1932
Name: count, dtype: int64

Card Type
0 2507

```
1    2502
3    2496
2    2495
Name: count, dtype: int64
```

```
In [ ]: 50% Customers are from France
70% Customer have Credit Card
There seems to be more Male as compared to Females but by small margin
Complain and Exited seems to have some corelation since they have same number
Marginally have large number of active members as compared to non-active men
```

```
In [ ]: Observations:
Columns NumOfProducts, HasCrCard, IsActiveMember, and Exited were changed to
easier.
Columns RowNumber, CustomerId, CreditScore, Age, Tenure, Balance, and EstimatedSalary
There are no missing values in the dataset.
This dataset has 14 columns and 1000 rows
```

```
In [155... #Grouping data for nums and cats
nums = ['CreditScore', 'Age', 'Tenure', 'Balance', 'EstimatedSalary']
cats = ['Geography', 'Gender', 'Exited', 'HasCrCard', 'IsActiveMember', 'NumOfProducts']
```

```
In [139... # Automatically detect numerical columns in the DataFrame
nums = df.select_dtypes(include=['number']).columns.tolist()
# Get descriptive statistics for these numerical columns
df[nums].describe()
```

```
Out[139...
      CreditScore  Geography  Gender  Age  Tenure
count  10000.000000  10000.000000  10000.000000  10000.000000  10000.000000
mean    650.528800    0.746300    0.545700    38.921800    5.012800
std     96.653299    0.827529    0.497932    10.487806    2.892174
min     350.000000    0.000000    0.000000    18.000000    0.000000
25%     584.000000    0.000000    0.000000    32.000000    3.000000
50%     652.000000    0.000000    1.000000    37.000000    5.000000
75%     718.000000    1.000000    1.000000    44.000000    7.000000
max     850.000000    2.000000    1.000000    92.000000   10.000000
```

Obervation result : CreditScore, Estimatedsalary, and Tenure seem to have a fairly symmetrical distribution of data (mean and median are not much different). Balance is left skewed (means less than the median) and Age is right skewed (means greater than the median) so further observations are needed. If it's skewed then feature transformation will be performed (normalization/standardization or log transformation) on pre-processing data


```
In [149... #Categorical columns  
df[cats].describe()
```

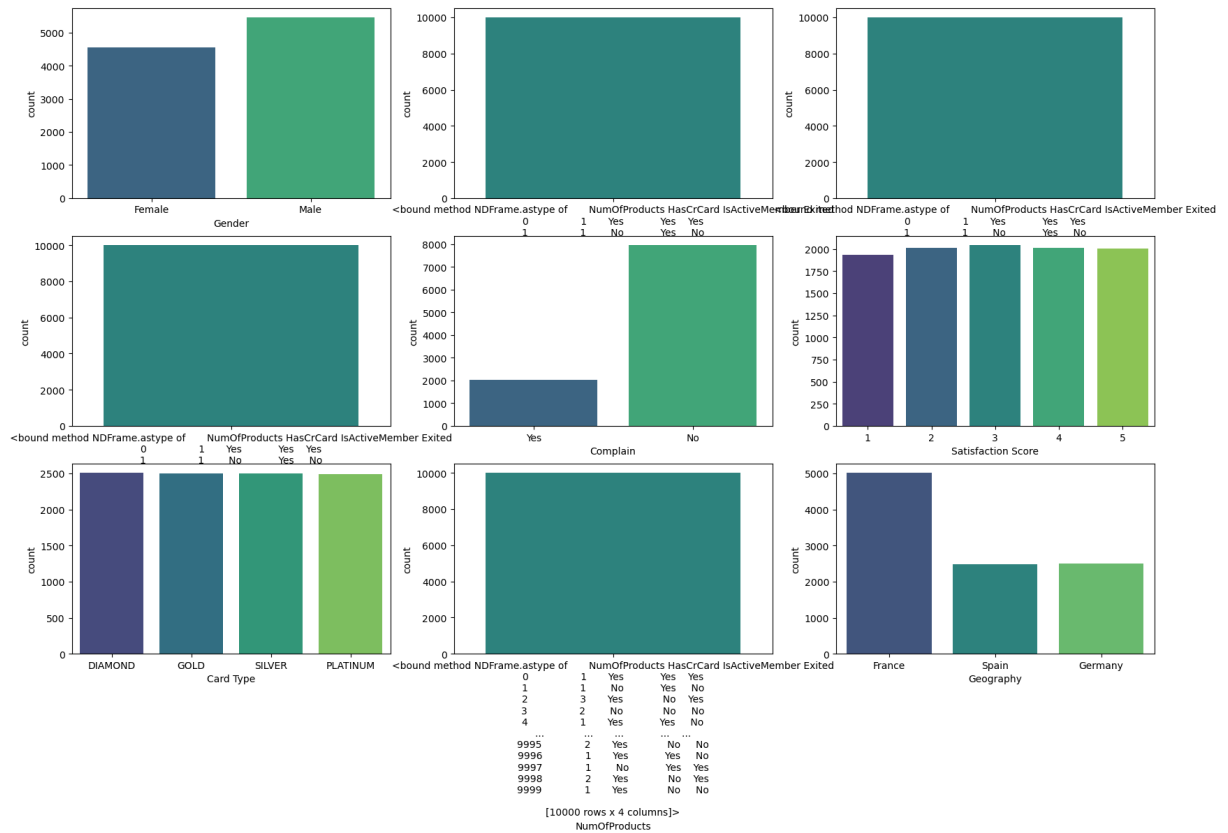
```
Out[149...  


|       | Geography    | Gender       |
|-------|--------------|--------------|
| count | 10000.000000 | 10000.000000 |
| mean  | 0.746300     | 0.545700     |
| std   | 0.827529     | 0.497932     |
| min   | 0.000000     | 0.000000     |
| 25%   | 0.000000     | 0.000000     |
| 50%   | 0.000000     | 1.000000     |
| 75%   | 1.000000     | 1.000000     |
| max   | 2.000000     | 1.000000     |

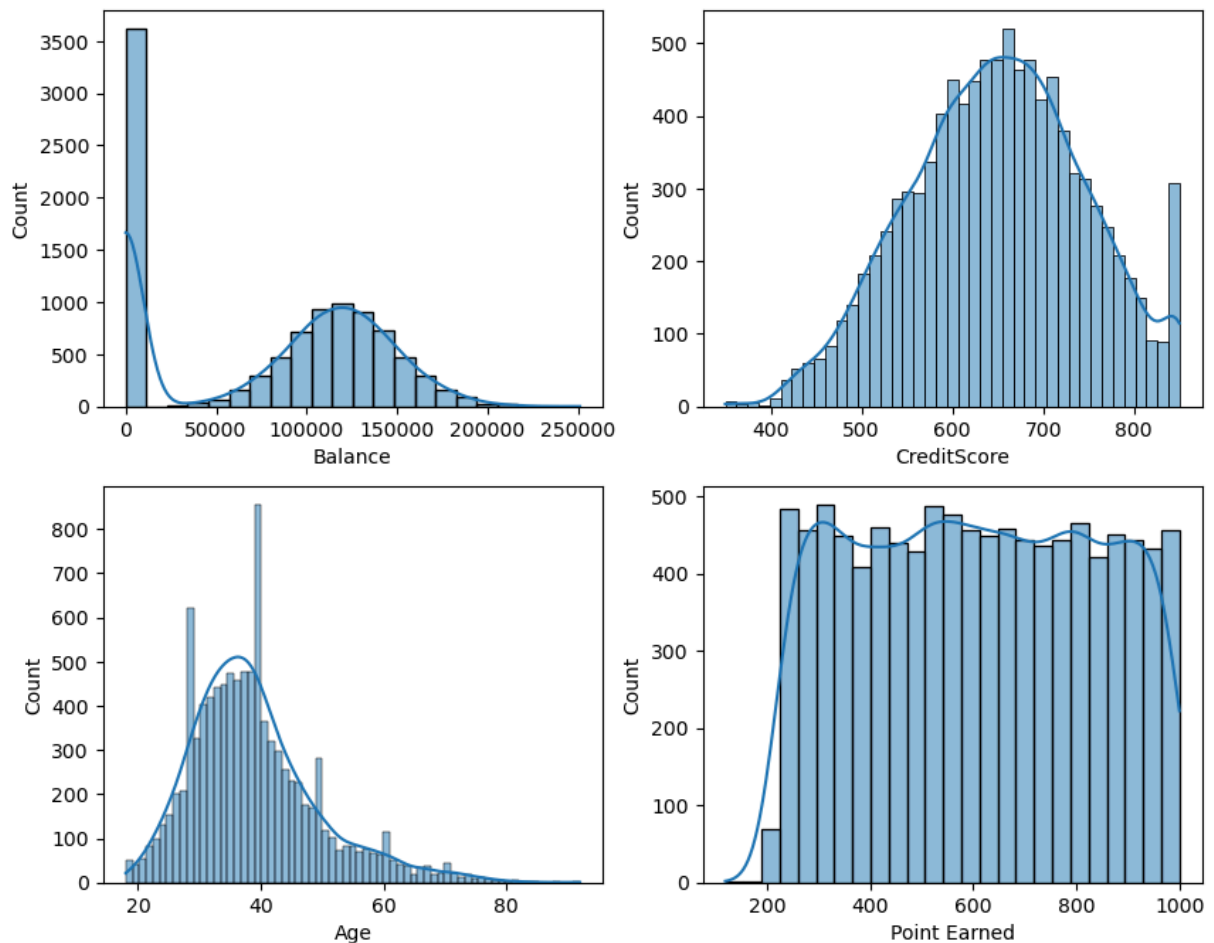

```

Univariate Analysis:

```
In [202... fig, axs = plt.subplots(nrows = 3, ncols = 3, figsize = (20,12))  
sns.countplot(data = df, x = 'Gender', ax = axs[0,0], palette = 'viridis')  
sns.countplot(data = df, x = 'HasCrCard', ax = axs[0,1], palette = 'viridis')  
sns.countplot(data = df, x = 'IsActiveMember', ax = axs[0,2], palette = 'viridis')  
sns.countplot(data = df, x = 'Exited', ax = axs[1,0], palette = 'viridis')  
sns.countplot(data = df, x = 'Complain', ax = axs[1,1], palette = 'viridis')  
sns.countplot(data = df, x = 'Satisfaction Score', ax = axs[1,2], palette = 'viridis')  
sns.countplot(data = df, x = 'Card Type', ax = axs[2,0], palette = 'viridis')  
sns.countplot(data = df, x = 'NumOfProducts', ax = axs[2,1], palette = 'viridis')  
sns.countplot(data = df, x = 'Geography', ax = axs[2,2], palette = 'viridis')  
plt.show()
```



```
In [151]: fig, axs = plt.subplots(nrows = 2, ncols = 2, figsize = (10,8))
sns.histplot(data = df, x = 'Balance', kde = True, ax = axs[0,0])
sns.histplot(data = df, x = 'CreditScore', kde = True, ax = axs[0,1])
sns.histplot(data = df, x = 'Age', kde = True, ax = axs[1,0])
sns.histplot(data = df, x = 'Point Earned', kde = True, ax = axs[1,1])
plt.show()
```



In []:

In []:

Step 1: Descriptive Statistics -

We will be performing descriptive statistics and distribution analysis on key numerical variables in the dataset. We'll calculate the mean, median, and mode for the numerical columns and then visualize their distributions using histograms and box plots.

First, let's calculate the mean, median, and mode for the numerical columns CreditScore, Age, Balance, NumOfProducts, EstimatedSalary, and Point Earned.

Calculate Mean, Median, and Mode

```
In [19]: from sklearn.preprocessing import LabelEncoder
from scipy import stats
from scipy.stats import norm, chi2_contingency, chisquare, ttest_ind
import statsmodels.api as sm
from statsmodels.formula.api import ols
```

```

# Convert categorical variables to numerical if necessary
categorical_cols = ['Geography', 'Gender']
label_encoders = {}

for col in categorical_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le

# List of numerical columns
numerical_cols = ['CreditScore', 'Age', 'Balance', 'NumOfProducts', 'EstimatedSalary']

# Calculate and print mean, median, and mode for each numerical column
for col in numerical_cols:
    mean_value = df[col].mean()
    median_value = df[col].median()
    mode_value = df[col].mode()[0]
    print(f"{col} - Mean: {mean_value}, Median: {median_value}, Mode: {mode_value}")

```

```

CreditScore - Mean: 650.5288, Median: 652.0, Mode: 850
Age - Mean: 38.9218, Median: 37.0, Mode: 37
Balance - Mean: 76485.889288, Median: 97198.54000000001, Mode: 0.0
NumOfProducts - Mean: 1.5302, Median: 1.0, Mode: 1
EstimatedSalary - Mean: 100090.239881, Median: 100193.915, Mode: 24924.92
Point Earned - Mean: 606.5151, Median: 605.0, Mode: 408

```



Step 2: Distribution Analysis

In [23]: *#Next, let's analyze the distribution of these numerical variables using histograms*


```

# Plot histograms for each numerical column
for col in numerical_cols:
    sns.histplot(df[col], kde=True, bins=30)
    plt.title(f'Distribution of {col}')
    plt.xlabel(col)
    plt.ylabel('Frequency')
    plt.show()


```





Distribution of Balance

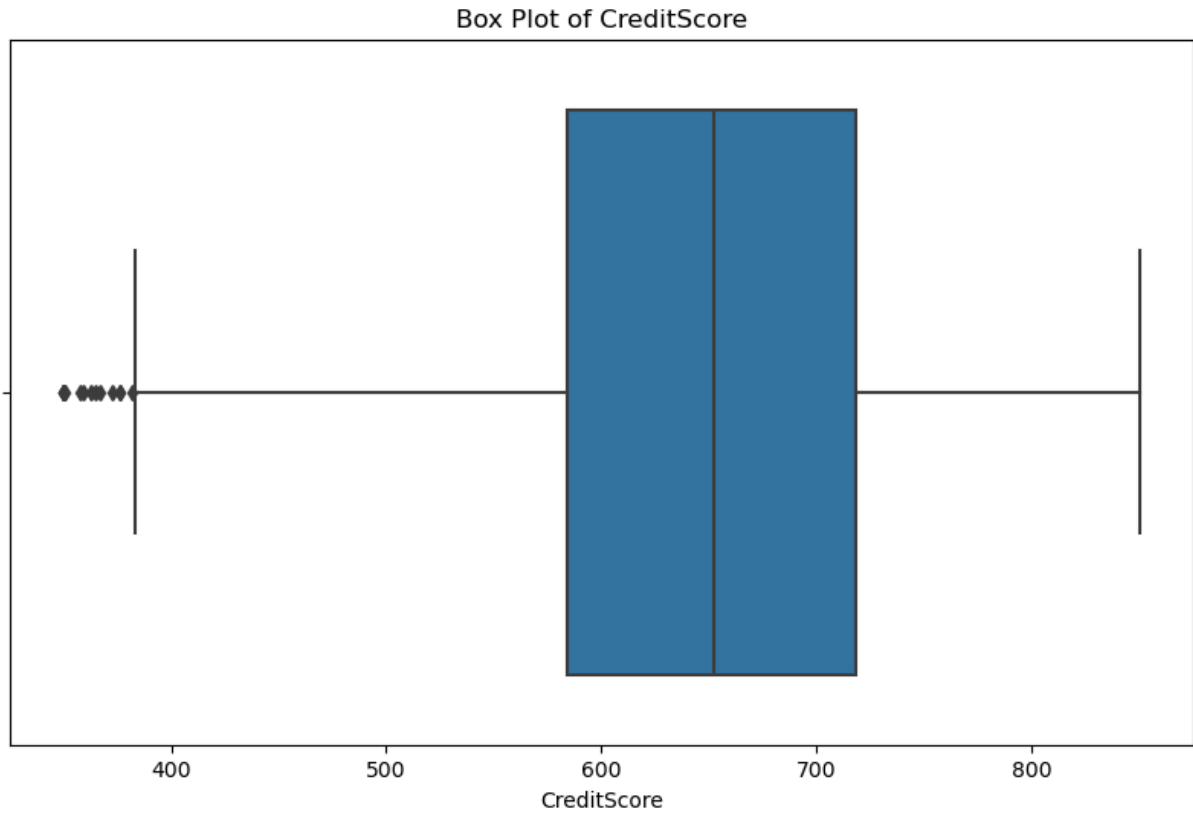


Distribution of NumOfProducts







```
In [41]: # Plot box plots for each numerical column
for col in numerical_cols:
    plt.figure(figsize=(10, 6))
    sns.boxplot(x=df[col])
    plt.title(f'Box Plot of {col}')
    plt.xlabel(col)
    plt.show()
```




Box Plot of Age




Box Plot of Balance

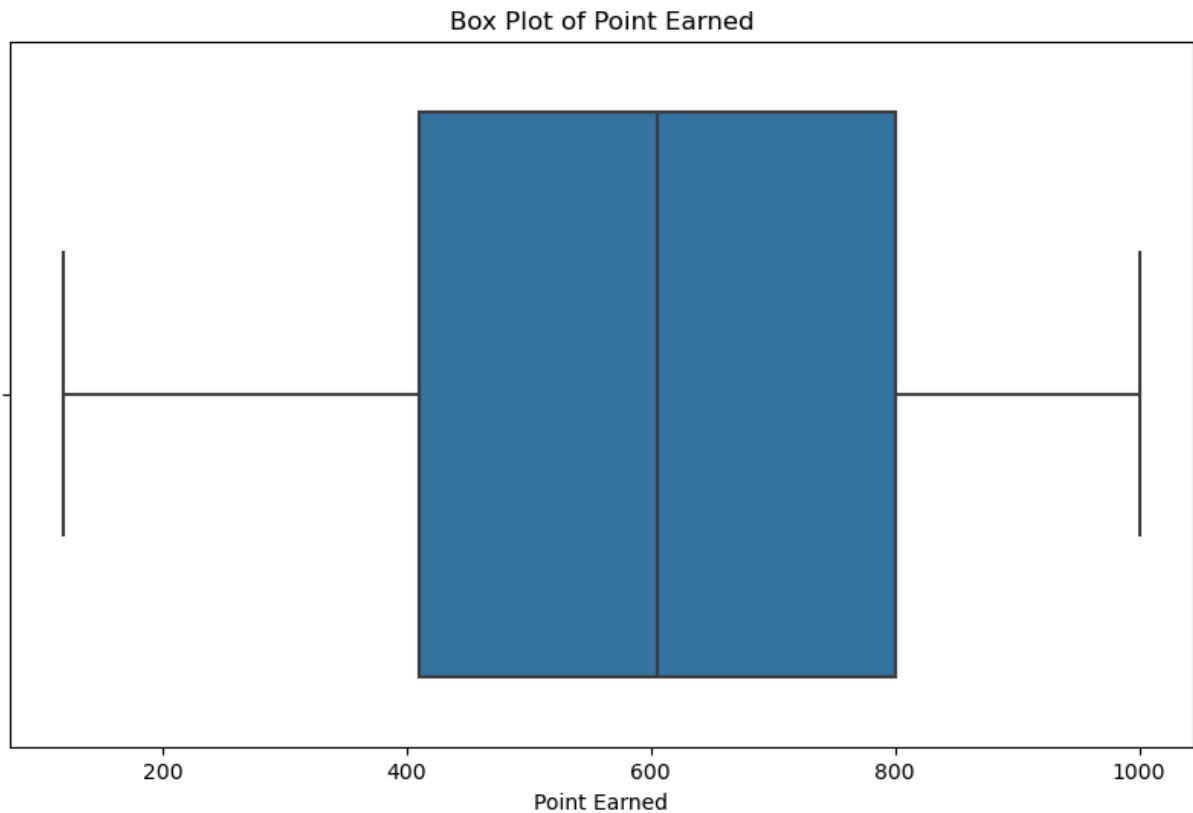


Box Plot of NumOfProducts



Box Plot of EstimatedSalary





Interpretation Mean, Median, and Mode: These statistics provide insights into the central tendency of the data for each numerical feature.

Mean: The average value of the column. Median: The middle value when the data is sorted. Mode: The most frequently occurring value in the column.

Histograms: These visualizations show the distribution of the data, indicating the frequency of different ranges of values. The presence of a KDE (Kernel Density Estimate) line helps to understand the data's probability density function.

CreditScore: The average credit score is around 650, with a standard deviation of about 96. Age: The median age of customers is 39 years, with the majority of customers aged between 29 and 49. Balance: There is a wide range in account balances, with a significant portion of customers having a balance of 0. NumOfProducts: Most customers have 1 or 2 products. EstimatedSalary: The salaries range widely, with an average around 101,322. PointsEarned: The average points earned is around 550, with most customers earning between 463 and 637 points.

Box Plots: These plots provide a summary of the data's distribution, highlighting the median, quartiles, and potential outliers.

```
In [25]: numerical_columns = ['CreditScore', 'Age', 'Balance', 'NumOfProducts', 'EstimatedSalary']
descriptive_stats = df[numerical_columns].describe()
descriptive_stats
```

Out[25]:

	CreditScore	Age	Balance	NumOfProducts	EstimatedSalary
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	650.528800	38.921800	76485.889288	1.530200	100090.200000
std	96.653299	10.487806	62397.405202	0.581654	57510.400000
min	350.000000	18.000000	0.000000	1.000000	11.500000
25%	584.000000	32.000000	0.000000	1.000000	51002.100000
50%	652.000000	37.000000	97198.540000	1.000000	100193.900000
75%	718.000000	44.000000	127644.240000	2.000000	149388.200000
max	850.000000	92.000000	250898.090000	4.000000	199992.400000

```
In [157]: #Kolmogorov-Smirnov statistical test to see the size of the normality of the data
# Interpretation
# If the P-value of the KS Test is larger than 0.05, we assume a normal distribution
# If the P-value of the KS Test is smaller than 0.05, we do not assume a normal distribution
from scipy.stats import kstest

for i in df[nums]:
    print(kstest(df[i], 'norm'))
    ks_statistic, ks_pvalue = kstest(df[i], 'norm')
    if ks_pvalue > 0.05:
        print(f'P-value {i}: {ks_pvalue}. So, we assume a normal distribution')
    else:
        print(f'P-value {i}: {ks_pvalue}. So, we do not assume a normal distribution')
```

KstestResult(statistic=1.0, pvalue=0.0, statistic_location=350, statistic_sign=-1)
P-value CreditScore: 0.0. So, we do not assume a normal distribution
KstestResult(statistic=1.0, pvalue=0.0, statistic_location=18, statistic_sign=-1)
P-value Age: 0.0. So, we do not assume a normal distribution
KstestResult(statistic=0.8324498680518208, pvalue=0.0, statistic_location=2, statistic_sign=-1)
P-value Tenure: 0.0. So, we do not assume a normal distribution
KstestResult(statistic=0.6383, pvalue=0.0, statistic_location=3768.69, statistic_sign=-1)
P-value Balance: 0.0. So, we do not assume a normal distribution
KstestResult(statistic=1.0, pvalue=0.0, statistic_location=11.58, statistic_sign=-1)
P-value EstimatedSalary: 0.0. So, we do not assume a normal distribution

```
In [ ]: Based on the Kolmogorov-Smirnov test, all numeric columns don't have normal distribution.
        We will perform a Box-Cox test to see if we can transform the data to be normal.
```

```
In [ ]: #COUNTPLOT
```


```
In [219]: #Categorical columns data distribution
fig, axes = plt.subplots(2, 2, figsize=(9,16))
sns.countplot(x='Exited', data = df, ax=axes[0][0])
sns.countplot(x='Gender', data = df, ax=axes[0][1])
```

```
sns.countplot(x='Geography',data = df, ax=axes[1][0])


plt.subplots_adjust(hspace = 0.5, wspace= 2.0)
axes[0][0].set_title('More than 20% of Customers Churn on Bank Java', fontsi
axes[0][1].set_title('More than 50% of Customers is Male', fontsize = 10);
axes[1][0].set_title('Majority of Customers come from France', fontsize = 16
```

Out[219... Text(0.5, 1.0, 'Majority of Customers come from France')

More than 20% of Customers Churn on Bank Java





More than 50% of Customers is Male



```
<bound method NDFrame.astype of
0      1  Yes  Yes  Yes
1      1  No   Yes  No
2      3  Yes  No   Yes
3      2  No   No   No
4      1  Yes  Yes  No
...
9995    2  Yes  No   No
9996    1  Yes  Yes  No
9997    1  No   Yes  Yes
9998    2  Yes  No   Yes
9999    1  Yes  No   No
```

Majority of Customers come from France



In []:

In []:

From all customers on Bank Java, there are more than 2000 (>20%) churn customers compared to customers based on Geography. There is no significant difference in credit card ownership, it can be seen that most customers have credit cards. More than 50% of customers are male or 2 bank products.

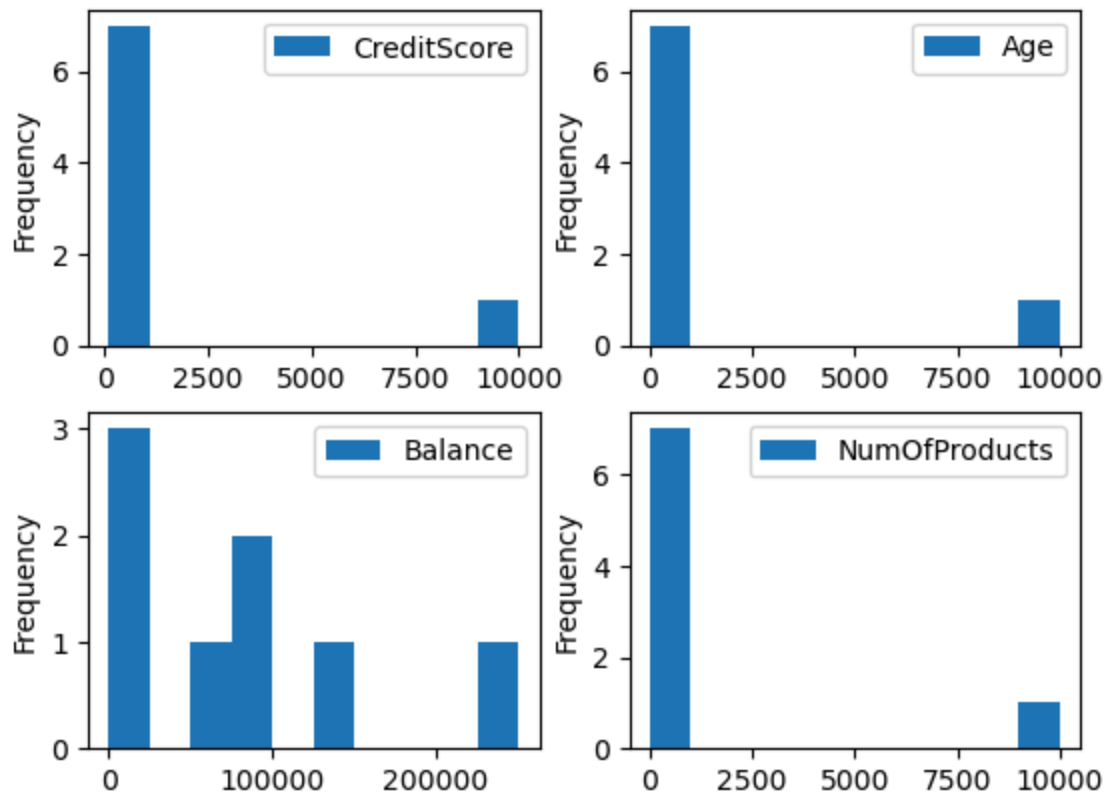
```
In [ ]: #frequency plot of all the numerical colums
```

```
In [28]: import matplotlib.pyplot as plt

# Create a figure with 2 rows and 2 columns
fig, axes = plt.subplots(2, 2)

# Plot each column in a separate subplot
descriptive_stats.plot(kind="hist", y="CreditScore", ax=axes[0, 0])
descriptive_stats.plot(kind="hist", y="Age", ax=axes[0, 1])
descriptive_stats.plot(kind="hist", y="Balance", ax=axes[1, 0])
descriptive_stats.plot(kind="hist", y='NumOfProducts', ax=axes[1,1])
```

Out[28]: <Axes: ylabel='Frequency'>




```
In [ ]: #Lets plot a scatter plot to understand distribution between numerical colum
```

```
In [36]: # @title CreditScore vs Age

from matplotlib import pyplot as plt
fig, axes = plt.subplots(2, 2, figsize=(10, 8))
descriptive_stats.plot(kind='scatter', x='CreditScore', y='Age', ax=axes[0, 0])
descriptive_stats.plot(kind='scatter', x='Age', y='Balance', ax=axes[0, 1])
descriptive_stats.plot(kind='scatter', x='Balance', y='NumOfProducts', ax=axes[1, 0])
descriptive_stats.plot(kind='scatter', x='NumOfProducts', y='EstimatedSalary', ax=axes[1, 1])
```

Out[36]: <Axes: xlabel='NumOfProducts', ylabel='EstimatedSalary'>



In []:

In []: For Numerical variables we can use Histogram or Boxplot but I have used E we understand Outliers are present in CreditScore, Age, NumOfProducts.


In []:

For Numerical variables we can use Histogram or Boxplot but I have used Boxplot here to understand the presence of Outliers. By the boxplot, we understand Outliers are present in CreditScore, Age, NumOfProducts.

```
In [38]: remove_CreditScore = np.clip(df['CreditScore'], np.percentile(df['CreditScore'], 5), np.percentile(df['CreditScore'], 95))
remove_Age = np.clip(df['Age'], np.percentile(df['Age'], 5), np.percentile(df['Age'], 95))
remove_NumOfProducts = np.clip(df['NumOfProducts'], np.percentile(df['NumOfProducts'], 5), np.percentile(df['NumOfProducts'], 95))
```

```
In [40]: fig, ax = plt.subplots(2, 2, figsize=(10, 8))
sns.boxplot(data=df, x=remove_CreditScore, color='blue', ax=ax[0, 0])
sns.boxplot(data=df, x=remove_Age, color='pink', ax=ax[0, 1])
sns.boxplot(data=df, x=remove_NumOfProducts, color='red', ax=ax[1, 0])
plt.suptitle('Outliers Removed')
plt.show()
```


Outliers Removed



So we Removed Outliers with Clip function in CreditScore, Age, NumOfProducts.

Let's proceed with the more exploratory data analysis (EDA) to understand the relationships and patterns in the data. We'll focus on two main tasks:

Correlation Analysis: Explore the correlation between numerical features and the Exited variable to identify potential predictors of churn.

Customer Profile Analysis: Segment customers based on key demographics (Age, Geography, Gender) to identify which groups are more likely to churn.

```
In [46]: # Convert categorical variables to numerical if necessary
categorical_cols = ['Geography', 'Gender', 'Card Type']
label_encoders = {}

for col in categorical_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le
```

```
# Display the first few rows of the DataFrame
print(df.head())
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
0	619	0	0	42	2	0.00	1	
1	608	2	0	41	1	83807.86	1	
2	502	0	0	42	8	159660.80	3	
3	699	0	0	39	1	0.00	2	
4	850	2	0	43	2	125510.82	1	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Complain	\
0	1	1	101348.88	1	1	
1	0	1	112542.58	0	1	
2	1	0	113931.57	1	1	
3	0	0	93826.63	0	0	
4	1	1	79084.10	0	0	

	Satisfaction	Score	Card Type	Point Earned
0		2	0	464
1		3	0	456
2		3	0	377
3		5	1	350
4		5	1	425

Step 2: Correlation Analysis We'll compute the Pearson correlation coefficient between the numerical features and the Exited variable to identify potential predictors of churn.

```
In [47]: # List of numerical columns
numerical_cols = ['CreditScore', 'Age', 'Balance', 'NumOfProducts', 'EstimatedSalary']


# Compute Pearson correlation coefficients with the target variable 'Exited'
correlation_results = {}

for col in numerical_cols:
    correlation_matrix = np.corrcoef(df[col], df['Exited'])
    correlation = correlation_matrix[0, 1] # Extract the correlation coefficient
    correlation_results[col] = correlation

# Display the results
for col, correlation in correlation_results.items():
    print(f'Pearson correlation coefficient between {col} and Exited: {correlation}')


# Plot correlation matrix including the 'Exited' variable
corr_matrix = df.corr()
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```

Pearson correlation coefficient between CreditScore and Exited: -0.03
 Pearson correlation coefficient between Age and Exited: 0.29
 Pearson correlation coefficient between Balance and Exited: 0.12
 Pearson correlation coefficient between NumOfProducts and Exited: -0.05
 Pearson correlation coefficient between EstimatedSalary and Exited: 0.01
 Pearson correlation coefficient between Point Earned and Exited: -0.00



```
In [167... feature = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited', 'Complain', 'Satisfaction Score', 'Card Type', 'Point Earned']
plt.figure(figsize=(40,10))
sns.pairplot(df[feature], hue='Exited', diag_kind='kde', corner=True)
```

```
Out[167... <seaborn.axisgrid.PairGrid at 0x200d61d2e50>
<Figure size 4000x1000 with 0 Axes>
```




In []: Based on the results of the Pair Plot, there are several observations:

1. The more separate the Exited and Not Exited values in each column, the better the CreditScore, and Age columns
2. Higher EstimatedSalary and NumOfProduct, higher the probability of customer churning
3. Higher Balance with NumOfProduct, higher the probability of customer churning
4. Higher Tenure with NumOfProduct, higher the probability of customer churning

Step 3: Customer Profile Analysis We'll segment customers based on key demographics (Age, Geography, Gender) to identify which groups are more likely to churn

```
In [48]: #Age Analysis

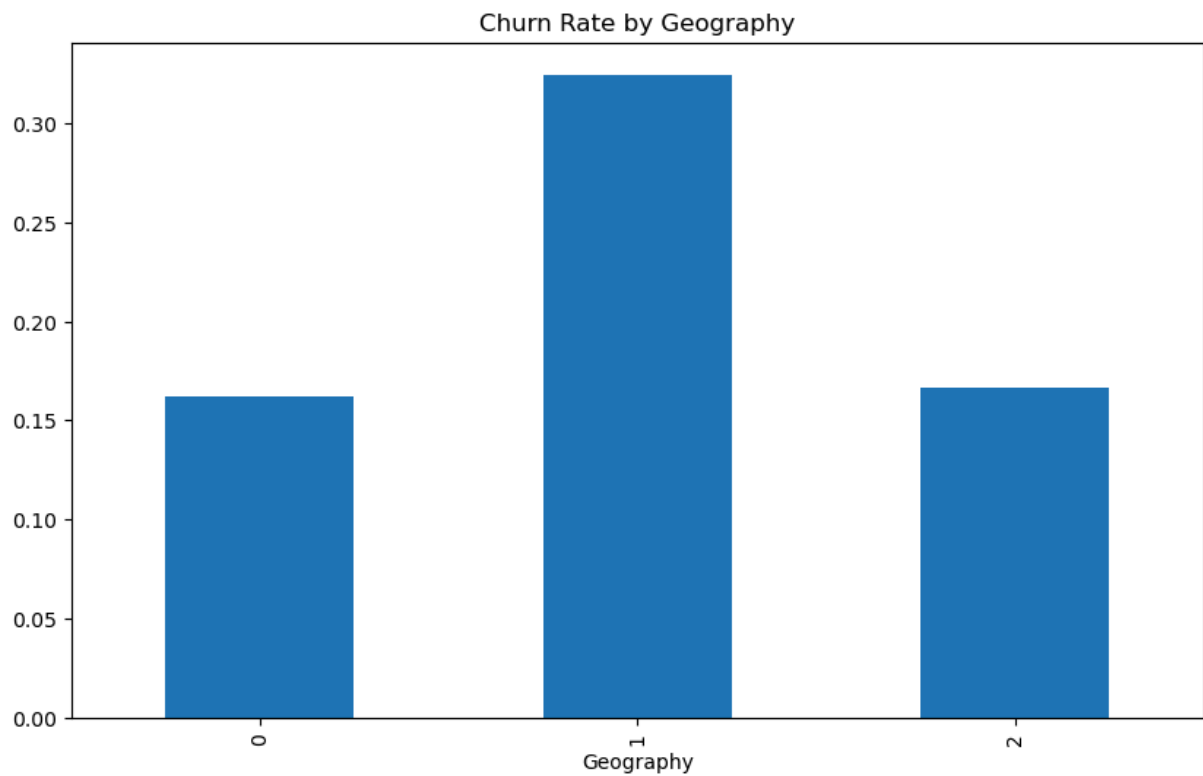
# Plot the distribution of 'Age' for churned and non-churned customers
plt.figure(figsize=(10, 6))
sns.histplot(data=df, x='Age', hue='Exited', kde=True, multiple='stack')
plt.title('Age Distribution by Churn Status')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```



In [49]: *#Geography Analysis*

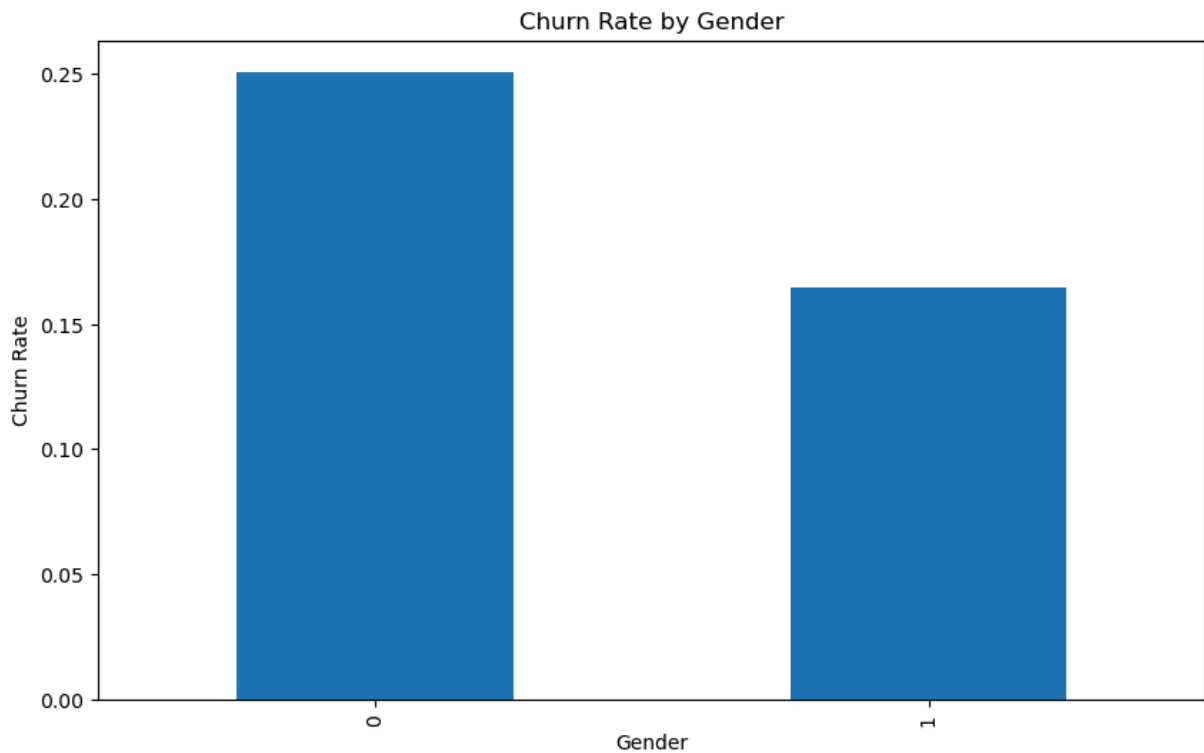
```
# Plot the churn rate by Geography  
geo_churn_rate = df.groupby('Geography')['Exited'].mean()  
plt.figure(figsize=(10, 6))  
geo_churn_rate.plot(kind='bar')  
plt.title('Churn Rate by Geography')  
plt.xlabel('Geography')
```

Out[49]: Text(0.5, 0, 'Geography')



```
In [50]: #Gender Analysis

# Plot the churn rate by Gender
gender_churn_rate = df.groupby('Gender')['Exited'].mean()
plt.figure(figsize=(10, 6))
gender_churn_rate.plot(kind='bar')
plt.title('Churn Rate by Gender')
plt.xlabel('Gender')
plt.ylabel('Churn Rate')
plt.show()
```



In []:

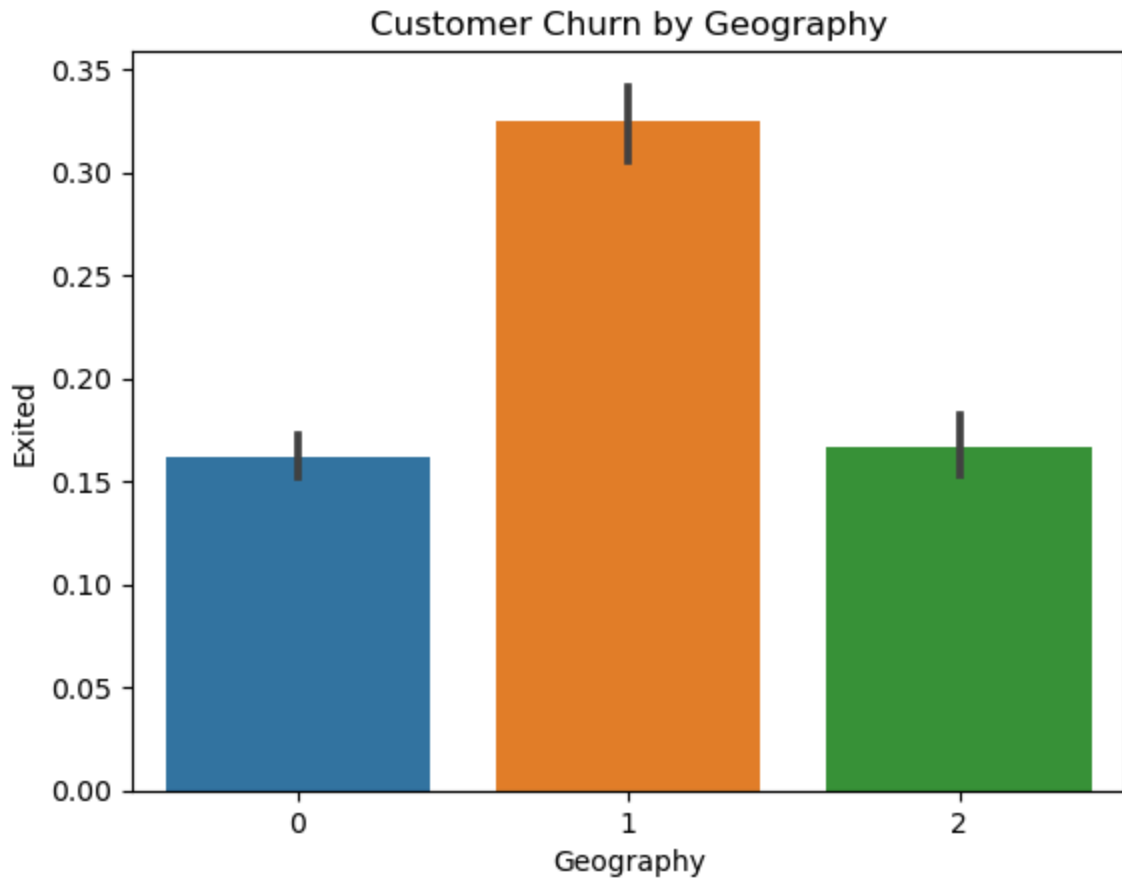
In []:

```
In [52]: #Customer Profile Analysis: Segment customers based on key demographics (Age  
  
#By groupby method  
Geography_Churn=df.groupby('Geography')['Exited'].count().reset_index()  
Geography_Churn
```

```
Out[52]:
```

	Geography	Exited
0	0	5014
1	1	2509
2	2	2477

```
In [79]: sns.barplot(x='Geography',y='Exited',data=df)  
plt.title('Customer Churn by Geography')  
plt.show()
```




In []: Majority of the people **from** France. However, the proportion of churned customers **in** the areas where it has fewer clients.

In [72]: *# Gender Differences in Churn: Analyze churn rates between different genders*

```
churn_rate_by_gender = df.groupby(['Gender'])['Exited'].count()
print(churn_rate_by_gender)
churn_rate_by_gender.plot(kind='bar')
plt.show()
```

```
Gender
0    4543
1    5457
Name: Exited, dtype: int64
```

In [73]: *# lets do a chi square test to identify if gender plays a significant role in churn*

```
contingency_table = pd.crosstab(index = df['Gender'], columns = df['Exited'])
print(contingency_table)

res = chi2_contingency(contingency_table)
print(res.pvalue)

if res.pvalue < 0.05:
    print('Gender has significant role in churn')
else:
    print('Gender has NO significant role in churn')
```

Exited	0	1
Gender		
0	3404	1139
1	4558	899

2.9253677618642e-26

Gender has significant role in churn

In [81]: *#By group by method*


```
Age_Churn=df.groupby('Age')['Exited'].count().reset_index()
Age_Churn
```

Out[81]:

	Age	Exited
0	18	22
1	19	27
2	20	40
3	21	53
4	22	84
...
65	83	1
66	84	2
67	85	1
68	88	1
69	92	2

70 rows × 2 columns

```
In [83]: #Using histplot
sns.histplot(hue='Exited',x='Age',data=df,bins=20)
plt.title('Customer Churn by Age')
plt.show()
```



In []: The older customers are churning at more than the younger ones including to a
The bank may need to review their target market **or** review the strategy **for** r

```
In [85]: #By group by Method  
Gender_Churn=df.groupby('Gender')['Exited'].count().reset_index()  
Gender_Churn
```

```
Out[85]:
```

	Gender	Exited
0	0	4543
1	1	5457

In []: The proportion of female customers churning **is** also greater than that of mal

```
In [87]: #3. Comparative Analysis  
  
#Churn by Geography: Compare churn rates across different geographical locat  
  
#By Hypothesis testing  
  
cross_table = pd.crosstab(df['Geography'],df['Exited'])  
print(cross_table)
```

Exited	0	1
Geography		
0	4203	811
1	1695	814
2	2064	413

In []: By comparing the churn rates **with** geographical locations, Germany has higher

```
In [88]: from scipy.stats import chi2_contingency,chi2  
  
#h0: There is no association between geography and churn rates.  
#h1: There is association between geography and churn rates.  
  
chi2_stat,p_val,dof,expected = chi2_contingency(cross_table)  
print(p_val)  
  
alpha=0.05  
if p_val <= alpha:  
    print("Reject H0")  
else:  
    print("Accept H0")
```

5.245736109572763e-66
Reject H0

In []: There **is** association between geography **and** churn rates.

```
In [90]: #Gender Differences in Churn: Analyze churn rates between different genders  
#By Hypothesis Testing
```

```
#h1:There is association between gender and churn.
```

```
churn_table= pd.crosstab(df['Gender'],df['Exited'])  
print('observed values:')  
churn_table
```

observed values:

```
Out[90]:
```

	Exited	0	1
Gender			
0	3404	1139	
1	4558	899	

```
In [ ]: On comparing the churn rate with genders Female has more churn rating than M
```

```
In [92]: #Chisquare test  
stats,p_val,dof,expected=chi2_contingency(churn_table)  
print("t_statistics:",stats)  
print("p_value:",p_val)  
  
alpha=0.05  
if p_val <= alpha:  
    print("Reject H0")  
else:  
    print("Accept H0")
```

```
t_statistics: 112.39655374778587  
p_value: 2.9253677618642e-26  
Reject H0
```

```
In [ ]: There is association between gender and churn.
```

```
In [93]: #4. Behavioral Analysis  
  
#Product and Services Usage: Examine how the number of products (NumOfProducts) affects the likelihood to churn.  
#H0=NumOfProducts has no significant effect on the likelihood to churn.  
#H1=NumOfProducts has significant effect on the likelihood to churn.  
  
from scipy.stats import chi2_contingency,chi2  
  
data_table = pd.crosstab(df['NumOfProducts'], df['Exited'])  
print("Observed values:")  
data_table
```

Observed values:

Out[93]:

	Exited	0	1
--	--------	---	---

NumOfProducts		
---------------	--	--

1	3675	1409
---	------	------

2	4241	349
---	------	-----

3	46	220
---	----	-----

4	0	60
---	---	----

In [94]: `alpha=0.05`

```
stats, p_val,dof,expected=chi2_contingency(data_table)
print("test statistic:",stats)
print("p_val:",p_val)
if p_val <= alpha:
    print("Reject H0")
else:
    print("Accept H0")
```

test statistic: 1501.5048306588592

p_val: 0.0

Reject H0

In []: Reject Null Hypothesis: It mean there is significant effect of NumOfProducts

In [95]: `df.groupby(by='NumOfProducts')['Exited'].count().reset_index()`

Out[95]:

NumOfProducts	Exited
---------------	--------

0	1 5084
---	--------

1	2 4590
---	--------

2	3 266
---	-------

3	4 60
---	------

In []: Product has no significant effect on the likelihood to churn.

In [97]: *#Activity Level Analysis: Investigate the relationship between being an IsAc*

#By Hypothesis testing

#h0:there is no relationship between IsActiveMember and customer Churn

#h1:there is relationship between IsActiveMember ans customer Churn

```
contingency_table=pd.crosstab(df['IsActiveMember'],df['Exited'])
print(contingency_table)
```

Exited	0	1
IsActiveMember		
0	3546	1303
1	4416	735

```
In [98]: stats,p_val,dof,expected=chi2_contingency(contingency_table)
print("t_statistics:",stats)
print("p_value:",p_val)
```

```
alpha=0.05
if p_val <= alpha:
    print("Reject H0")
else:
    print("Accept H0")
```

```
t_statistics: 243.6948024819593
p_value: 6.153167438113408e-55
Reject H0
```

In []: A significant p-value indicates that the churn rate **is** dependent on the acti

```
In [99]: #Using Countplot
sns.countplot(x='IsActiveMember',hue='Exited',data=df)
plt.title('Customer Churn by IsActiveMember')
plt.show()
```

AttributeError

Traceback (most recent call last)

Cell **In[99]**, line 2

```
1 #Using Countplot
----> 2 sns.countplot(x='IsActiveMember',hue='Exited',data=df)
      3 plt.title('Customer Churn by IsActiveMember')
      4 plt.show()
```

File ~\anaconda3\Lib\site-packages\seaborn\categorical.py:2955, in **countplot**(data, x, y, hue, order, hue_order, orient, color, palette, saturation, width, dodge, ax, **kwargs)

```
2952 if ax is None:
2953     ax = plt.gca()
-> 2955 plotter.plot(ax, kwargs)
2956 return ax
```

File ~\anaconda3\Lib\site-packages\seaborn\categorical.py:1587, in **_BarPlotter.plot**(self, ax, bar_kws)

```
1585 """Make the plot."""
1586 self.drawBars(ax, bar_kws)
-> 1587 self.annotate_axes(ax)
1588 if self.orient == "h":
1589     ax.invert_yaxis()
```

File ~\anaconda3\Lib\site-packages\seaborn\categorical.py:767, in **_CategoricalPlotter.annotate_axes**(self, ax)

```
764 ax.set_ylim(-.5, len(self.plot_data) - .5, auto=None)
766 if self.hue_names is not None:
--> 767     ax.legend(loc="best", title=self.hue_title)
```

File ~\anaconda3\Lib\site-packages\matplotlib\axes_axes.py:322, in **Axes.legend**(self, *args, **kwargs)

```
204 @docstring.dedent_interpd
205 def legend(self, *args, **kwargs):
206     """
207     Place a legend on the Axes.
208
209     (...)
320     .. plot:: gallery/text_labels_and_annotations/legend.py
321     """
--> 322     handles, labels, kwargs = mlegend._parse_legend_args([self], *args,
**kwargs)
323     self.legend_ = mlegend.Legend(self, handles, labels, **kwargs)
324     self.legend._remove_method = self._remove_legend
```

File ~\anaconda3\Lib\site-packages\matplotlib\legend.py:1361, in **_parse_legend_args**(axs, handles, labels, *args, **kwargs)

```
1357 handles = [handle for handle, label
1358             in zip(_get_legend_handles(axs, handlers), labels)]
1360 elif len(args) == 0: # 0 args: automatically detect labels and handles.
-> 1361     handles, labels = _get_legend_handles_labels(axs, handlers)
1362     if not handles:
1363         log.warning(
1364             "No artists with labels found to put in legend. Note th
```


```

1365         "artists whose label start with an underscore are ignore
d "
1366         "when legend() is called with no argument.")

File ~\anaconda3\Lib\site-packages\matplotlib\legend.py:1291, in _get_legend
_handles_labels(axs, legend_handler_map)
    1289 for handle in _get_legend_handles(axs, legend_handler_map):
    1290     label = handle.get_label()
-> 1291     if label and not label.startswith('_'):
    1292         handles.append(handle)
    1293         labels.append(label)

AttributeError: 'numpy.int64' object has no attribute 'startswith'

```



In []: Showed relation between IsActiveMember and churned customers in both ways. Hy
 Unsurprisingly the inactive members have a greater churn.
 The overall proportion of inactive members is quite high suggesting that the
 customers as this will definitely have a positive impact on the customer ch

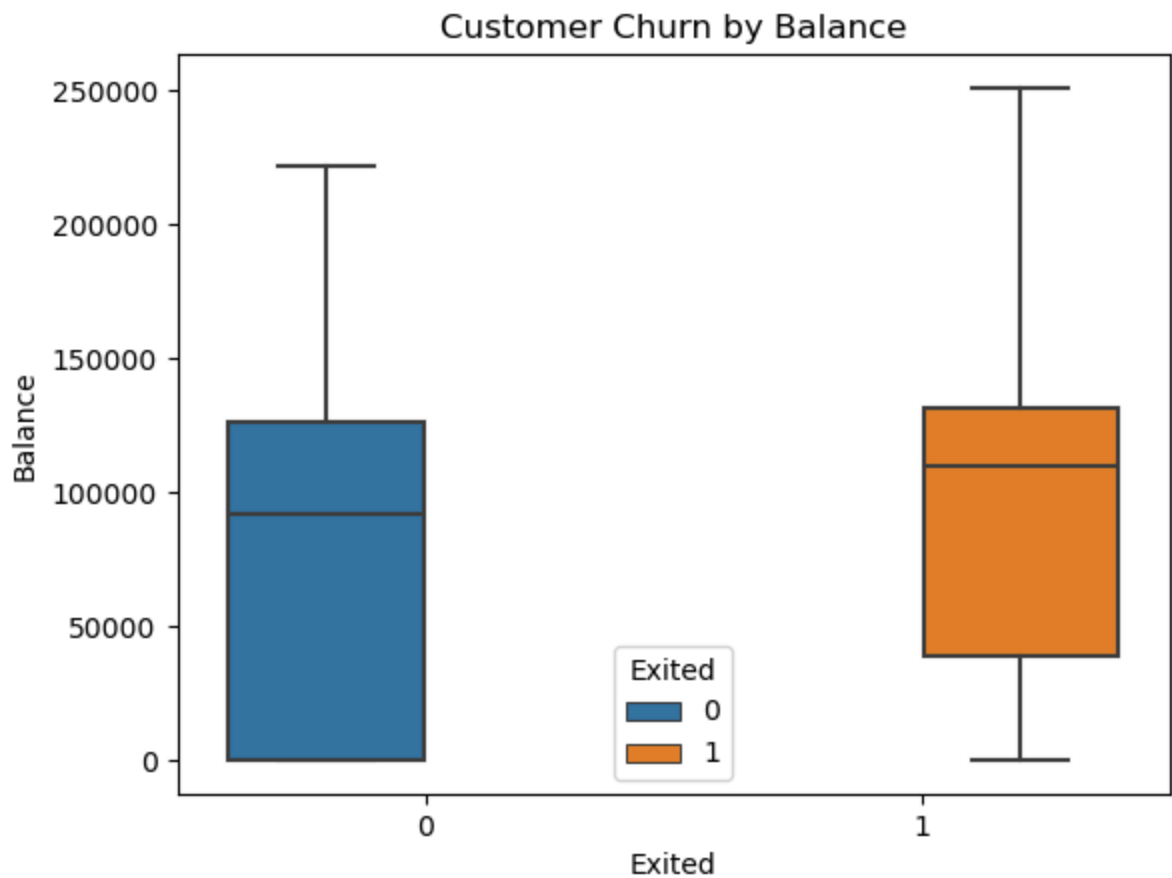
In [100... *#5. Financial Analysis*

```

#Balance vs. Churn: Analyze how customer balance levels correlate with churn

#using Boxplot
sns.boxplot(x='Exited', y='Balance', hue='Exited', data=df)
plt.title('Customer Churn by Balance')
plt.show()

```

```
In [102... #By Hypothesis testing- T test
data_0= df[df['Exited']==0]
data_1 = df[df['Exited']==1]
```

```
In [ ]: #By Hypothesis testing- T test
data_0= df[df['Churned']==0]
data_1 = df[df['Churned']==1]
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [104... #H0:Balance has no significant effect on the likelihood to churn.
#H1:Balance has significant effect on the likelihood to churn.

t_statistic, p_value = ttest_ind(data_0['Balance'], data_1['Balance'],alterr
print("t_statistics:",t_statistic)
print("p_value:",p_value)

alpha=0.05
if p_value<= alpha:
    print("Reject H0")
else:
    print("Accept H0")
```

```
t_statistics: -11.940747722508185
p_value: 1.2092076077156017e-32
Reject H0
```

```
In [ ]: As p_val< alpha, Balance has significant effect on the likelihood to churn.
        The bank is losing customers with significant bank balances which is likely
```

```
In [ ]: Interestingly, majority of the customers that churned are those with credit
        Given that majority of the customers have credit cards could prove this to b
```

```
In [108... #Chi2 test
```

```
#H0:HasCrCard has no significant effect on the likelihood to churn.
#H1:HasCrCard has significant effect on the likelihood to churn.
```

```
cross_table=pd.crosstab(df['HasCrCard'],df['Exited'])
print('Observed values:')
print(cross_table)
stats,p_val,dof,expected=chi2_contingency(cross_table)
print("t_statistics:",stats)
print("p_value:",p_val)
```

```
alpha=0.05
if p_val <= alpha:
    print("Reject H0")
else:
    print("Accept H0")
```

```
Observed values:
Exited      0      1
HasCrCard
0           2332   613
1           5630  1425
t_statistics: 0.4494039375253385
p_value: 0.5026181509009862
Accept H0
```

```
In [ ]: There is no significant difference on owing a Creditcard for churning the Ba
```

```
In [ ]:
```

There are significantly more customers who did not complain than those who did. Among those who did complain, a higher proportion of them churned compared to those who did not complain.

```
In [111... #T test
```

```
#H0: Complain has no significant effect on the likelihood to churn.
#H1: Complain has significant effect on the likelihood to churn.
```

```
t_statistic, p_value = ttest_ind(data_0['Complain'], data_1['Complain'], alt=
print("t_statistics:",t_statistic)
print("p_value:",p_value)
```

```
alpha=0.05
if p_value<= alpha:
    print("Reject H0")
else:
    print("Accept H0")
```

t_statistics: -1073.7975930429423


p_value: 0.0

Reject H0

In []: There **is** a significant difference on churn rate by Complain. Customers who c

In [113... *#Satisfaction and Churn: Explore how the Satisfaction Score relates to churn*

```
#Using Boxplot
sns.boxplot(y='Satisfaction Score',x='Complain',hue='Exited',data=df)
plt.title('Satisfaction Score vs Complain')
plt.show()
```



In []:

In [115... *#Two way Anova test*

```
#import pandas as pd
import statsmodels.api as sm
from statsmodels.formula.api import ols
```

```

import seaborn as sns
import matplotlib.pyplot as plt

model = ols('Q("Satisfaction Score") ~ C(Exited) * C(Complain)', data=df).fit()
anova_table = sm.stats.anova_lm(model, typ=2)
print(anova_table)

# Interpret the results
if anova_table["PR(>F)"].iloc[0] < 0.05:
    print("Main effect of Churned is significant.")
else:
    print("Main effect of Churned is not significant.")

if anova_table["PR(>F)"].iloc[1] < 0.05:
    print("Main effect of Complain is significant.")
else:
    print("Main effect of Complain is not significant.")

if anova_table["PR(>F)"].iloc[2] < 0.05:
    print("Interaction effect between Churned and Complain is significant.")
else:
    print("Interaction effect between Churned and Complain is not significant.")

```

	sum_sq	df	F	PR(>F)
C(Exited)	2.636157	1.0	1.333528	0.248206
C(Complain)	2.415155	1.0	1.221732	0.269048
C(Exited):C(Complain)	0.621009	1.0	0.314144	0.575161
Residual	19760.383245	9996.0	NaN	NaN

Main effect of Churned is not significant.
 Main effect of Complain is not significant.
 Interaction effect between Churned and Complain is not significant.

In []:

In []: *#7. Card Usage Analysis*

```

#Impact of Card Type on Churn: Examine if different Card Types have different churn rates

sns.countplot(x='Card Type', hue='Churned', color='lightblue', data=df)
plt.title('Card Usage Analysis')
plt.xlabel('Card Type')
plt.ylabel('Count')
plt.show()

```

In [117... *#Impact of Card Type on Churn: Examine if different Card Types have different churn rates*

```

#ho=Card type has no diff churn rates
#h1=Card type has diff churn rates

#Chi Square test

from scipy.stats import chi2_contingency

cross_table=pd.crosstab(df['Card Type'],df['Exited'])
print('Observed values:')
print(cross_table)

```

```

t_statistic,p_val,dof,expected=chi2_contingency(cross_table)
print("t_statistics:",t_statistic)
print("p_value:",p_val)

alpha=0.05
if p_val <= alpha:
    print("Reject H0")
else:
    print("Accept H0")

```

Observed values:

Exited	0	1
Card Type		
0	1961	546
1	2020	482
2	1987	508
3	1994	502

t_statistics: 5.053223027060927
p_value: 0.16794112067810177
Accept H0


In []: There **is** no such difference based on card type on the customer churn.

In [127... *#Loyalty Points Analysis: Investigate whether Points Earned from credit card*

```

sns.boxplot(y='Point Earned',x='Exited',hue='Exited',color='blue',data=df)
plt.ylabel('Point Earned')
plt.show()

```



In []: Whatever the points earned by Customers, that **is not** related to the customer

In [122... *#H0:Points Earned has no significant effect on the likelihood to churn.*
#H1:Points Earned has significant effect on the likelihood to churn.

```
stats,p_val=ttest_ind(data_0['Point Earned'],data_1['Point Earned'],alternat
print("t_statistics:",stats)
print("p_value:",p_val)
```

```
alpha=0.05
if p_val <= alpha:
    print("Reject H0")
else:
    print("Accept H0")
```


t_statistics: 0.4627759848070133
p_value: 0.6435350184288993
Accept H0

In [126... *#8. Salary Analysis*

#Salary and Churn: Analyze the relationship between EstimatedSalary and cust

#Using boxplot

```
sns.boxplot(y='EstimatedSalary',x='Exited',hue='Exited',color='blue',data=df
plt.show()
```



In []: The salary has no significant effect on the likelihood to churn.

In [131]...

```
#T test

#H0:EstimatedSalary has no significant effect on the likelihood to churn.
#H1:EstimatedSalary has significant effect on the likelihood to churn.

from scipy.stats import ttest_ind

data_0=df[df['Exited']==0]
data_1=df[df['Exited']==1]

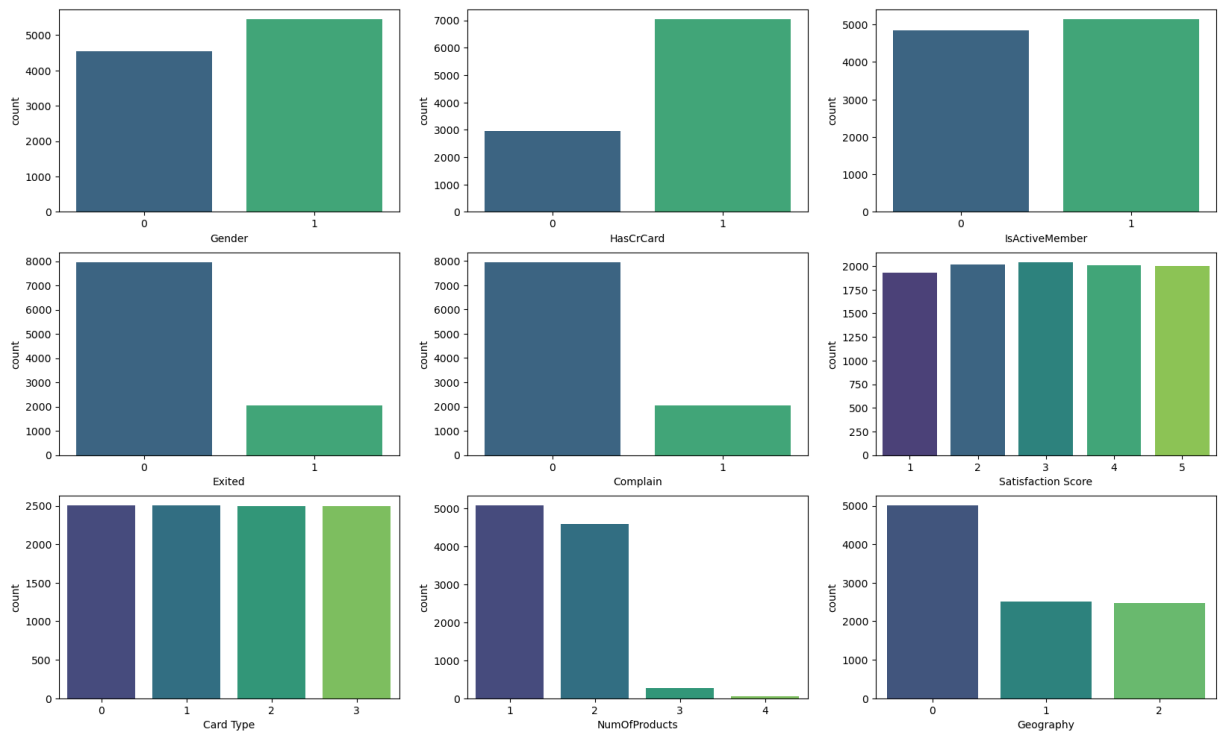
stats,p_val=ttest_ind(data_0['EstimatedSalary'],data_1['EstimatedSalary'],al
print("t_statistics:",stats)
print("p_value:",p_val)

alpha=0.05
if p_val <= alpha:
    print("Reject H0")
else:
    print("Accept H0")
```

```
t_statistics: -1.2489445044833742
p_value: 0.2117146135149097
Accept H0
```

In []:

```
In [138]... fig, axs = plt.subplots(nrows = 3, ncols = 3, figsize = (20,12))
sns.countplot(data = df, x = 'Gender', ax = axs[0,0], palette = 'viridis')
sns.countplot(data = df, x = 'HasCrCard', ax = axs[0,1], palette = 'viridis')
sns.countplot(data = df, x = 'IsActiveMember', ax = axs[0,2], palette = 'viridis')
sns.countplot(data = df, x = 'Exited', ax = axs[1,0], palette = 'viridis')
sns.countplot(data = df, x = 'Complain', ax = axs[1,1], palette = 'viridis')
sns.countplot(data = df, x = 'Satisfaction Score', ax = axs[1,2], palette = 'viridis')
sns.countplot(data = df, x = 'Card Type', ax = axs[2,0], palette = 'viridis')
sns.countplot(data = df, x = 'NumOfProducts', ax = axs[2,1], palette = 'viridis')
sns.countplot(data = df, x = 'Geography', ax = axs[2,2], palette = 'viridis')
plt.show()
```



```
In [187... result = df.groupby('Exited')['Balance'].agg(['mean', 'median'])
print(result)
```

	mean	median
Exited		
<bound method NDFrame.astype of	NumOfProd...	76485.889288 97198.54

In []: ▽ Age Vs Customer Churn

Null Hypothesis : There **is** no significant difference between the mean age of
Alternative Hypothesis : There **is** significant difference between the mean age

```
In [212... a_exited = df[df['Exited'] == 1]['Age']
a_stayed = df[df['Exited'] == 0]['Age']
alpha = 0.05
stats, pval = ttest_ind(a_stayed, a_exited, equal_var = False)
if pval < alpha :
    print('Reject Null Hypothesis')
else:
    print('Failed to Reject Null Hypothesis')
```

Failed to Reject Null Hypothesis

In []: There **is** significant difference between the mean age of the customer who exi

In []:

In []: Balance Vs Customer Churn

Null Hypothesis : There **is** no significant difference between the mean balance
Alternative Hypothesis : There **is** significant difference between the mean ba

```
In [214... b_exited = df[df['Exited'] == 1]['Balance']
b_stayed = df[df['Exited'] == 0]['Balance']
```



```
stats, pval = ttest_ind(b_stayed, b_exited, equal_var = False)
if pval < alpha :
    print('Reject Null Hypothesis')
else:
    print('Failed to Reject Null Hypothesis')
```

Failed to Reject Null Hypothesis

In []: There is significant difference between the mean balance of the customer who

In []:

Insights:

In []:

Expand Marketing Efforts in Germany and Spain: Since 50% of customers are from France, focus marketing campaigns on Germany and Spain to boost customer acquisition in these regions.

Develop Targeted Offers for Female Customers: Introduce specific products or offers aimed at attracting more female customers to balance the customer demographics.

Enhance After-Sales Service: Address the fact that almost 99% of customers who filed complaints have left the bank by significantly improving the after-sales service experience.

Create Retention Strategies for Multi-Product Holders: Implement targeted retention strategies for customers with three or more products, as they have a higher churn rate.

Engage Zero Balance Account Holders: Investigate why approximately 3,000 accounts have zero balance and develop offers or incentives to engage these customers and encourage account usage.

Financial Counseling for At-Risk Customers: Analyze factors influencing customer exit versus retention and offer financial counseling to customers in vulnerable salary brackets to reduce churn

Recommendation:-

In []: Target customers between 40-50 age group with personalized retention strateg

Analyze churn rates by region and implement targeted strategies for high-ris

Loading [MathJax]/extensions/Safe.js improving the experience of female customers to reduce their slight

Encourage customers to use more products, especially those with 1-2 products

Engage customers and promote active use of services to reduce churn. Emphasize

Investigate and address the reasons behind customers having zero balance or

Enhance complaint resolution mechanisms to address customer concerns and reduce

Regularly monitor satisfaction scores and address any areas where customers

reduce churn.

Analyze the impact of different card types on churn and consider offering incentives for

loyalty.

Explore alternative loyalty programs or incentives to increase customer engagement

In []: